

# covid19\_project

November 10, 2025

Adam Glogovský and Oleksii Bieliaiev

## 1 COVID-19 Outcome Modeling

### 1.0.1 All Imports

```
[ ]: import warnings
from copy import deepcopy
from pathlib import Path

import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
import wandb
from IPython.display import display
from sklearn.metrics import (
    accuracy_score,
    average_precision_score,
    classification_report,
    confusion_matrix,
    f1_score,
    precision_recall_curve,
    precision_score,
    recall_score,
    roc_auc_score,
)
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from torch.utils.data import DataLoader, TensorDataset, WeightedRandomSampler

warnings.filterwarnings("ignore")
plt.style.use("seaborn-v0_8")
torch.manual_seed(42)
np.random.seed(42)
```

Define path to dataset, batch size for training and device (cuda for NVIDIA)

```
[ ]: DATA_PATH = Path("dataset.csv")
BATCH_SIZE = 1024
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Device: {device}")
```

Load and explore dataset: show shape, numeric columns, distributions, and correlation heatmap

```
[ ]: raw_df = pd.read_csv(DATA_PATH)
print(f"Shape of dataset: {raw_df.shape}")
display(raw_df.head())

numeric_df = raw_df.select_dtypes(include=["number"])

raw_df.info()
display(raw_df.count(numeric_only=True).to_frame(name="count"))
print("Dataset shape:", raw_df.shape)

for col in numeric_df.columns:
    series = numeric_df[col]
    counts = series.value_counts(dropna=False).sort_index()
    x = range(len(counts))
    labels = counts.index.astype(str)
    plt.figure(figsize=(8, 3))
    plt.bar(x, counts.values, color="steelblue", alpha=0.9, width=0.8)
    plt.title(f"Count per value - {col}", fontsize=12, pad=8)
    plt.xlabel("Value")
    plt.ylabel("Count")
    if len(labels) > 40:
        step = max(1, len(labels) // 40)
        plt.xticks(list(x)[::step], labels[::step], rotation=45, ha="right")
    else:
        plt.xticks(x, labels, rotation=45, ha="right")
    plt.grid(True, linestyle="--", axis="y", alpha=0.4)
    plt.tight_layout()
    plt.show()

corr = numeric_df.corr()
plt.figure(figsize=(10, 10))
sns.heatmap(
    corr,
    annot=True,
    cmap="coolwarm",
    fmt=".1f",
    square=True,
    linewidths=1,
    cbar_kws={"shrink": 0.7, "label": "Correlation"})
```

```
)
plt.title("Correlation Heatmap of Numeric Attributes", fontsize=14, pad=12)
plt.tight_layout()
plt.show()
```

## 1.1 Exploratory analysis

```
[ ]: display(raw_df.describe(include='all').transpose())
```

Create binary target (1 = COVID, 0 = No COVID) and visualize target and age distributions

```
[ ]: target_series = raw_df['CLASIFICATION_FINAL'].isin([1, 2, 3]).astype(int)
fig, axes = plt.subplots(1, 2, figsize=(10, 4))
target_series.value_counts(normalize=True).mul(100).reindex([0, 1]).plot(
    kind='bar', ax=axes[0], color=['steelblue', 'tomato'])
)
axes[0].set_title('Target distribution (%)')
axes[0].set_xlabel('HAS_COVID')
axes[0].set_ylabel('Percent')

age_sample = raw_df['AGE'].clip(0, 100)
axes[1].hist(age_sample, bins=30, color='slategray', alpha=0.9)
axes[1].set_title('Age distribution (trimmed 0-100)')
axes[1].set_xlabel('Age')
axes[1].set_ylabel('Count')
plt.tight_layout()
plt.show()
```

```
[ ]: numeric_cols = raw_df.select_dtypes(include=['number']).columns
sample_for_corr = raw_df[numeric_cols].sample(n=min(2000, len(raw_df)), ↴
    random_state=42)
plt.figure(figsize=(10, 8))
sns.heatmap(sample_for_corr.corr(), cmap='coolwarm', linewidths=0.1, center=0)
plt.title('Correlation heatmap (raw numeric sample)')
plt.tight_layout()
plt.show()
```

```
[ ]: demo_df = raw_df.copy()
print('Sample before preprocessing:')
display(demo_df.iloc[:30])

demo_df.columns = demo_df.columns.str.strip()
demo_df.rename(columns={'DATE_DIED': 'DEAD'}, inplace=True)
med_unit_original = demo_df['MEDICAL_UNIT'].copy()
classif_original = demo_df['CLASIFICATION_FINAL'].copy()

demo_df['DEAD'] = np.where(demo_df['DEAD'] == '9999-99-99', 2, 1)
```

```

demo_df.loc[demo_df['SEX'] == 1, 'PREGNANT'] = 0
binary_cols = [
    'USMER', 'SEX', 'INTUBED', 'PNEUMONIA', 'PREGNANT', 'DIABETES', 'COPD', 'ASTHMA', 'INMSUPR', 'HIPER',
    'CLASIFICATION_FINAL'
]
demo_df[binary_cols] = demo_df[binary_cols].replace(2, 0)
demo_df[binary_cols] = demo_df[binary_cols].replace(97, 0.5)
demo_df[binary_cols] = demo_df[binary_cols].replace(98, 0.5)
demo_df['MEDICAL_UNIT'] = med_unit_original
demo_df['CLASIFICATION_FINAL'] = classif_original
demo_df.rename(columns={'SEX': 'IS_MALE'}, inplace=True)
demo_df['CLASIFICATION_FINAL'] = demo_df['CLASIFICATION_FINAL'].\
    replace([1,2,3], 1)
demo_df['CLASIFICATION_FINAL'] = demo_df['CLASIFICATION_FINAL'].\
    replace([4,5,6,7], 0)
demo_df.rename(columns={'CLASIFICATION_FINAL': 'HAS_COVID'}, inplace=True)
print('Sample after key preprocessing steps:')
display(demo_df.iloc[:30])
# Display purpose only

```

## 1.2 Preprocessing

Our preprocessing function simply prepares the dataframe we have, for model. We work with only valid values, retype non-numeric values to numeric, normalize data by leaving 1 (true) be 1, but 2 (false) to be 0 and unknown values with np.nan. We ensure that male patients cannot have pregnant value set to 1. # Most importantly we replace values in final classification from 1-6 to be only 1 or 2 redefining it to binary classification. (it now represents if the patient is sick, or is not sick). And we also normalize the age of patients

```

[ ]: # Preprocessing function that we actually use
def preprocess_data(frame: pd.DataFrame) -> pd.DataFrame:
    df = frame.copy() # work on a copy to avoid modifying the original
    df.rename(columns={"SEX": "IS_MALE"}, inplace=True)

    # Clean column names and rename gender column for clarity
    df.columns = df.columns.str.strip()
    df.rename(columns={"SEX": "IS_MALE"}, inplace=True)

    # Keep only rows with valid COVID classification values (1-6)
    valid = df["CLASIFICATION_FINAL"].isin([1, 2, 3, 4, 5, 6])
    df = df.loc[valid].copy()

    # Create binary column DEAD (1 = dead, 0 = alive)
    df["DEAD"] = np.where(df["DATE_DIED"] == "9999-99-99", 0, 1)

    # List of binary features (0/1 type)
    binary_cols = [
        "USMER", "IS_MALE", "INTUBED", "PNEUMONIA", "PREGNANT", "DIABETES",
        "COPD",
    ]

```

```

    "ASTHMA", "INMSUPR", "HIPERTENSION", "OTHER_DISEASE", "CARDIOVASCULAR",
    "OBESITY", "RENAL_CHRONIC", "TOBACCO", "ICU", "DEAD",
]

# Replace invalid codes (97-99 etc.) with NaN to mark missing data
replace_map = {1: 1, 2: 0, 3: np.nan, 97: np.nan, 98: np.nan, 99: np.nan}
df[binary_cols] = df[binary_cols].replace(replace_map)

# Ensure males are never marked as pregnant
df.loc[df["IS_MALE"] == 1, "PREGNANT"] = 0

# Create hospitalization flag: 1 = hospitalized, 0 = outpatient
df["IS_HOSPITALIZED"] = df["PATIENT_TYPE"].map({1: 0, 2: 1}).fillna(0)
df.drop(columns=["PATIENT_TYPE"], inplace=True)

# Convert COVID classification (1-3 positive, 4-6 negative) into binary
# label
df["HAS_COVID"] = df["CLASIFICATION_FINAL"].replace({1: 1, 2: 1, 3: 1, 4: 0,
                                                       5: 0, 6: 0})
df.drop(columns=["CLASIFICATION_FINAL", "DATE_DIED"], inplace=True)

# Fill missing binary values with median (usually 0 or 1)
df[binary_cols] = df[binary_cols].fillna(df[binary_cols].median())

# Define chronic condition columns
risk_cols = [
    "DIABETES", "COPD", "ASTHMA", "INMSUPR", "HIPERTENSION",
    "OTHER_DISEASE", "CARDIOVASCULAR", "OBESITY", "RENAL_CHRONIC",
]

# Create aggregated health indicators
df["CHRONIC_COUNT"] = df[risk_cols].sum(axis=1) # number of
# chronic diseases
df["MULTI_MORBID"] = (df["CHRONIC_COUNT"] >= 2).astype(int) # has 2+
# chronic diseases
df["RISK_OBESE_SMOKER"] = ((df["OBESITY"] == 1) | (df["TOBACCO"] == 1)).\
    astype(int) # obese or smoker

# Bucketize age into groups (0-30, 31-45, 46-60, 61-75, 76+)
df["AGE_BUCKET"] = (
    pd.cut(df["AGE"], bins=[0, 30, 45, 60, 75, 120], labels=False,
           include_lowest=True)
    .fillna(0)
    .astype(int)
)

```

```

# Normalize age using standard scaling (mean=0, std=1)
scaler = StandardScaler()
df["AGE"] = scaler.fit_transform(df[["AGE"]])

# One-hot encode medical unit column (convert category → binary columns)
df = pd.get_dummies(df, columns=["MEDICAL_UNIT"], prefix="MED_UNIT", ↴
drop_first=True)

# Drop any remaining missing values and reset index
df = df.dropna()
df.reset_index(drop=True, inplace=True)

# Return cleaned and fully preprocessed dataset
return df

```

```
[ ]: model_df = preprocess_data(raw_df)
print(f"Processed shape: {model_df.shape}")
display(model_df.head())
feature_cols = [col for col in model_df.columns if col != "HAS_COVID"]
print(f"Feature count: {len(feature_cols)}")
```

```
[ ]: processed_numeric = model_df.select_dtypes(include=['number'])
corr_processed = processed_numeric.corr()
plt.figure(figsize=(10, 10))
sns.heatmap(
    corr_processed,
    annot=True,
    cmap='coolwarm',
    fmt='.1f',
    square=True,
    linewidths=1,
    cbar_kws={'shrink': 0.7, 'label': 'Correlation'}
)
plt.title('Correlation Heatmap after Preprocessing', fontsize=14, pad=12)
plt.tight_layout()
plt.show()
```

### 1.3 Data split

Next we split data for training (70%) for validation (10%) and for testing (20%) so we can accurately train our NN

```
[ ]: # Split dataset into train (70%), validation (10%), and test (20%) sets
train_df, temp_df = train_test_split(
    model_df,
    test_size=0.3,                                     # 30% goes to temp (val + test)
```

```

    stratify=model_df["HAS_COVID"],                      # keep same class ratio across splits
    random_state=42,                                         # ensure reproducibility
    shuffle=True,
)

# Split temp into validation (1/3 of temp = 10%) and test (2/3 of temp = 20%)
val_df, test_df = train_test_split(
    temp_df,
    test_size=2 / 3,
    stratify=temp_df["HAS_COVID"],
    random_state=42,
    shuffle=True,
)

# Reset indices for clean dataframes
train_df = train_df.reset_index(drop=True)
val_df = val_df.reset_index(drop=True)
test_df = test_df.reset_index(drop=True)

# Helper function to print class distribution info
def describe_split(name, frame):
    share = frame["HAS_COVID"].mean() * 100
    print(f"{name}: {len(frame)} | positive={share:.2f}% | negative={100 - share:.2f}%")

# Display dataset split summary
print("Dataset splits:")
describe_split("Train", train_df)
describe_split("Val", val_df)
describe_split("Test", test_df)

```

## 1.4 Torch datasets

From now on we will work with PyTorch, so we need to convert our DataFrames to TensorDatasets

```
[ ]: # Convert DataFrame to PyTorch TensorDataset (features + labels)
def frame_to_dataset(frame):
    features = frame[feature_cols].to_numpy(dtype=np.float32)      # input features
    labels = frame["HAS_COVID"].to_numpy(dtype=np.float32)          # target labels
    return TensorDataset(torch.from_numpy(features), torch.from_numpy(labels))

# Create datasets for training, validation, and testing
train_dataset = frame_to_dataset(train_df)
val_dataset = frame_to_dataset(val_df)
test_dataset = frame_to_dataset(test_df)
```

```

# Compute class weights to handle imbalance (more weight for minority class)
class_counts = train_df["HAS_COVID"].value_counts().to_dict()
weights = train_df["HAS_COVID"].map(lambda value: 1.0 / class_counts[value]).  

    ↪to_numpy(dtype=np.float32)

# Create a sampler for balanced batches during training
sampler = WeightedRandomSampler(weights=torch.from_numpy(weights),  

    ↪num_samples=len(weights), replacement=True)

# Build data loaders for efficient batching
train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, sampler=sampler)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE)

# Calculate positive class weight for loss function (to handle imbalance)
positive = float((train_df["HAS_COVID"] == 1).sum())
negative = float((train_df["HAS_COVID"] == 0).sum())
pos_weight_value = negative / max(positive, 1.0)

# Define binary classification loss with class weight
criterion = nn.BCEWithLogitsLoss(pos_weight=torch.tensor([pos_weight_value]),  

    ↪dtype=torch.float32, device=device)

# Inspect one batch from the training loader
sample_batch = next(iter(train_loader))
print(f"Sample batch → X: {sample_batch[0].shape}, y: {sample_batch[1].shape},  

    ↪pos_weight={pos_weight_value:.2f}")

```

## 1.5 Model and helpers

Using PyTorch integrated linear function, activation function and other cool things, we create our model specialised for binary classification, predicting if the patient has Covid or not. # We also calculate metrics like f1 score to determine how good is our NN trained.

```
[ ]: # Simple feedforward MLP for binary COVID classification
class CovidMLP(nn.Module):
    def __init__(self, input_dim, hidden_layers, dropout):
        super().__init__()
        layers = []
        prev = input_dim # start with the input feature size

        # build hidden layers based on config
        for hidden in hidden_layers:
            layers.append(nn.Linear(prev, hidden)) # dense layer
            layers.append(nn.BatchNorm1d(hidden)) # helps stabilize
            ↪training
            layers.append(nn.ReLU()) # non-linearity
```

```

        layers.append(nn.Dropout(dropout))           # regularization
        prev = hidden                                # next layer input =_
        ↵this layer output

        # final output layer (1 neuron for binary classification)
        layers.append(nn.Linear(prev, 1))
        self.net = nn.Sequential(*layers)

    def forward(self, x):
        # forward pass through the network
        return self.net(x).squeeze(1)   # remove extra dimension (batch, 1) →_
        ↵(batch,)

# Calculate common metrics from probabilities
def metrics_from_probs(probs, labels, threshold=0.5):
    preds = (probs >= threshold).astype(int)  # turn probs into 0/1 predictions

    metrics = {
        "accuracy": accuracy_score(labels, preds),
        "precision": precision_score(labels, preds, zero_division=0),
        "recall": recall_score(labels, preds, zero_division=0),
        "f1": f1_score(labels, preds, zero_division=0),
        "pr_auc": average_precision_score(labels, probs),  # area under_
        ↵precision-recall curve
    }

    # how much recall we get when precision is at least 0.7
    prec_curve, rec_curve, _ = precision_recall_curve(labels, probs)
    mask = prec_curve >= 0.7
    metrics["recall_at_70_precision"] = float(rec_curve[mask].max()) if mask.any() else 0.0

    return metrics

# Run evaluation on val/test data and collect metrics
def evaluate_model(model, loader):
    model.eval()  # disable dropout, switch to eval mode
    losses, probs, labels = [], [], []

    with torch.no_grad():  # no gradients needed for evaluation
        for xb, yb in loader:
            xb = xb.to(device)
            yb = yb.to(device)

            logits = model(xb)          # forward pass

```

```

        loss = criterion(logits, yb)    # compute loss
        losses.append(loss.item())

        # convert logits → probabilities and collect labels
        probs.append(torch.sigmoid(logits).cpu())
        labels.append(yb.cpu())

# merge all batches into full arrays
prob_array = torch.cat(probs).numpy()
label_array = torch.cat(labels).numpy()

# compute metrics
metrics = metrics_from_probs(prob_array, label_array)
metrics["roc_auc"] = roc_auc_score(label_array, prob_array)
metrics["loss"] = float(np.mean(losses)) # average val loss

return metrics, prob_array, label_array

# Find the threshold that gives the best F1 score
def find_best_threshold(probs, labels):
    best_thr = 0.5
    best_score = 0.0

    # try thresholds from 0.1 to 0.9
    for thr in np.linspace(0.1, 0.9, 17):
        score = f1_score(labels, (probs >= thr).astype(int), zero_division=0)
        if score > best_score:
            best_score = score
            best_thr = float(thr)

    return best_thr, best_score

```

Our model training function builds the model from the config, initializes the optimizer, scheduler, and logs everything using Wandb. # Then it runs set amount of epochs from config, training the model, changing the learning rate and logging more data, and saving the best result for later

```
[ ]: def train_model(config):
    # Build model from config and move to device
    model = CovidMLP(len(feature_cols), config["hidden_layers"], ▾
    config["dropout"]).to(device)

    # Optimizer: AdamW (Adam + decoupled weight decay for better generalization)
    optimizer = torch.optim.AdamW(
        model.parameters(),
        lr=config["lr"],
        weight_decay=config["weight_decay"]
```

```

)

# LR scheduler: reduce LR when validation loss plateaus (mode='min')
# factor=0.5 → halve LR; patience=1 → wait 1 epoch without improvement
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer, mode="min", factor=0.5, patience=1
)

# Initialize Weights & Biases run for experiment tracking
# - project: group runs
# - config: hyperparams logged as run config
# - reinit=False: reuse the same process/run context
wandb_run = wandb.init(project="covid-19", config=config, reinit=False)
wandb_run.watch(model, log="all") # log gradients/weights for debugging

best_state = None # best model weights (by val F1)
best_snapshot = None # (metrics, probs, labels) for best epoch
best_score = -np.inf # track best F1
history = [] # per-epoch metrics for plotting

# ---- Training loop ----
for epoch in range(1, config["epochs"] + 1):
    model.train() # enable dropout/batchnorm updates
    epoch_losses = []

    # Iterate mini-batches
    for xb, yb in train_loader:
        xb = xb.to(device)
        yb = yb.to(device)

        optimizer.zero_grad() # reset accumulated grads
        logits = model(xb) # forward pass → logits
        loss = criterion(logits, yb) # BCEWithLogitsLoss (with_
        ↪pos_weight)
        loss.backward() # backpropagate gradients
        optimizer.step() # update weights

    epoch_losses.append(loss.item())

# ---- Validation ----
val_metrics, val_probs, val_labels = evaluate_model(model, val_loader)

# Step LR scheduler with validation loss (expects a scalar)
scheduler.step(val_metrics["loss"])

# Save epoch summary for later analysis/plots
history.append({

```

```

    "epoch": epoch,
    "train_loss": float(np.mean(epoch_losses)),
    **{k: v for k, v in val_metrics.items()},
})

# Console log
print(
    f"{config['name']:<10s} epoch {epoch:02d} "
    f"train_loss={history[-1]['train_loss']:.3f} "
    f"val_loss={val_metrics['loss']:.3f} "
    f"val_f1={val_metrics['f1']:.3f}"
)

# Log key metrics to W&B for dashboards/curves
wandb_run.log({
    "epoch": epoch,
    "train_loss": float(np.mean(epoch_losses)),
    "val_loss": val_metrics["loss"],
    "val_f1": val_metrics["f1"],
    "learning_rate": optimizer.param_groups[0]["lr"],
})

# Track best model by validation F1
if val_metrics["f1"] > best_score:
    best_score = val_metrics["f1"]
    best_state = deepcopy(model.state_dict()) # keep ↵weights
    best_snapshot = (val_metrics.copy(), val_probs.copy(), # keep ↵artifacts
                     val_labels.copy())

# Restore best weights before returning
model.load_state_dict(best_state)

# Close W&B run cleanly
wandb_run.finish()

# Return trained (best) model, full training history, and best-epoch ↵snapshot
return model, history, best_snapshot

```

## 1.6 Experiments

We define a few experiments, with different names so we can track them later. Each of these experiments is a config for our model training function, after each training we save the results of the experiment so we can compare them later.

```
[ ]: # Define a small grid of experiments (architectures + hyperparameters)
experiment_grid = [
    {"name": "baseline", "hidden_layers": [128, 64, 32], "dropout": 0.3, "lr": 3e-4, "weight_decay": 1e-4, "epochs": 10},
    {"name": "compact", "hidden_layers": [64, 32], "dropout": 0.25, "lr": 5e-4, "weight_decay": 5e-5, "epochs": 12},
    {"name": "wide", "hidden_layers": [256, 128, 64], "dropout": 0.4, "lr": 2e-4, "weight_decay": 1e-4, "epochs": 12},
]
experiment_results = [] # will store trained model snapshots + metrics for each config

# Run training for each config in the grid
for cfg in experiment_grid:
    model, history, snapshot = train_model(cfg) # train and return best-epoch snapshot
    metrics_at_05, val_probs, val_labels = snapshot # metrics at 0.5 threshold & raw probs/labels

    # Tune decision threshold to maximize F1 on validation set
    best_thr, _ = find_best_threshold(val_probs, val_labels)

    # Recompute metrics using the tuned threshold (keep ROC AUC & loss from snapshot)
    tuned_metrics = metrics_from_probs(val_probs, val_labels, threshold=best_thr)
    tuned_metrics["roc_auc"] = metrics_at_05["roc_auc"]
    tuned_metrics["loss"] = metrics_at_05["loss"]

    # Save experiment artifacts for later comparison / model selection
    experiment_results.append(
        {
            "name": cfg["name"], # experiment label
            "config": cfg, # full hyperparameter
            "state_dict": deepcopy(model.state_dict()), # best model weights
            "threshold": best_thr, # tuned decision threshold
            "val_metrics": tuned_metrics, # metrics at tuned threshold
            "history": history, # per-epoch training/val
        }
    )

# Build a compact comparison table across experiments
```

```

results_table = pd.DataFrame(
    [
        {
            "experiment": result["name"],
            "val_f1": result["val_metrics"]["f1"],
            "val_pr_auc": result["val_metrics"]["pr_auc"],
            "val_recall_at_70_precision": result["val_metrics"]["recall_at_70_precision"],
            "threshold": result["threshold"],
        }
        for result in experiment_results
    ]
).sort_values(by="val_f1", ascending=False) # rank by F1

print("Validation summary:")
display(results_table.reset_index(drop=True))

```

```

[ ]: # Plot training history for each experiment (train loss, val loss, F1 across epochs)
history_metrics = ["train_loss", "loss", "f1"]

for result in experiment_results:
    history_df = pd.DataFrame(result["history"]) # convert per-epoch metrics to DataFrame
    plt.figure(figsize=(7, 4))

    # Plot selected metrics if present
    for metric in history_metrics:
        if metric in history_df.columns:
            plt.plot(history_df["epoch"], history_df[metric], label=metric.replace("_", " "))

    plt.title(f"Training history - {result['name']}") # experiment name
    plt.xlabel("Epoch")
    plt.ylabel("Value")
    plt.grid(alpha=0.3)
    plt.legend()
    plt.tight_layout()
    plt.show()

```

## 1.7 Test evaluation

Now we pick the best model from experiments (based on F1 score), evaluate the model and save the metrics, create classification report, confusion matrix and log it to Wandb

```
[ ]:
```

```

# Select best experiment by highest validation F1
best_experiment = max(experiment_results, key=lambda item:_
    item["val_metrics"]["f1"])
best_cfg = best_experiment["config"] # best hyperparameters
best_threshold = best_experiment["threshold"] # tuned decision threshold (from_
    ↪val set)

# Recreate the best model architecture and load its saved weights
best_model = CovidMLP(len(feature_cols), best_cfg["hidden_layers"],_
    ↪best_cfg["dropout"]).to(device)
best_model.load_state_dict(best_experiment["state_dict"])

# Evaluate on the held-out test set
test_metrics_raw, test_probs, test_labels = evaluate_model(best_model,_
    ↪test_loader)

# Recompute metrics with the tuned threshold (keep ROC AUC & loss from raw eval)
test_metrics = metrics_from_probs(test_probs, test_labels,_
    ↪threshold=best_threshold)
test_metrics["roc_auc"] = test_metrics_raw["roc_auc"]
test_metrics["loss"] = test_metrics_raw["loss"]

# Pretty-print main test metrics
print("Test metrics with tuned threshold:")
for key in ["loss", "roc_auc", "pr_auc", "accuracy", "precision", "recall",_
    ↪"f1", "recall_at_70_precision"]:
    print(f" {key}: {test_metrics[key]:.4f}")

# Derive hard predictions using tuned threshold, then print a detailed_
    ↪classification report
binary_preds = (test_probs >= best_threshold).astype(int)
print("Classification report:")
print(classification_report(test_labels, binary_preds, digits=4,_
    ↪zero_division=0))

# Confusion matrix (as a DataFrame for nicer display)
cm = confusion_matrix(test_labels, binary_preds)
cm_df = pd.DataFrame(cm, index=["Actual 0", "Actual 1"], columns=["Pred 0",_
    ↪"Pred 1"])
print("Confusion matrix:")
display(cm_df)

# Log confusion matrix to Weights & Biases for the best config
wandb_run = wandb.init(project="covid-19", config=best_cfg, reinit=True) # new_
    ↪run for test artifacts
wandb_run.log({

```

```

    "confusion_matrix": wandb.plot.confusion_matrix(
        preds=binary_preds.tolist(),
        y_true=test_labels.tolist(),
        title="Confusion Matrix of best Model"
    )
})

wandb_run.finish()

# Plot Precision-Recall curve with the chosen operating point highlighted
precisions, recalls, _ = precision_recall_curve(test_labels, test_probs)
plt.figure(figsize=(6, 4))
plt.plot(recalls, precisions, label="PR curve")
plt.scatter(test_metrics["recall"], test_metrics["precision"], color="red", ▾
    label="Chosen threshold")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Test precision-recall curve")
plt.grid(alpha=0.3)
plt.legend()
plt.show()

```