

covid19_project_with_outputs

November 10, 2025

Adam Glogovský and Oleksii Bieliaiev

1 COVID-19 Outcome Modeling

1.0.1 All Imports

```
[1]: import warnings
from copy import deepcopy
from pathlib import Path

import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
import wandb
from IPython.display import display
from sklearn.metrics import (
    accuracy_score,
    average_precision_score,
    classification_report,
    confusion_matrix,
    f1_score,
    precision_recall_curve,
    precision_score,
    recall_score,
    roc_auc_score,
)
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from torch.utils.data import DataLoader, TensorDataset, WeightedRandomSampler

warnings.filterwarnings("ignore")
plt.style.use("seaborn-v0_8")
torch.manual_seed(42)
np.random.seed(42)
```

Define path to dataset, batch size for training and device (cuda for NVIDIA)

```
[2]: DATA_PATH = Path("dataset.csv")
BATCH_SIZE = 1024
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Device: {device}")
```

Device: cuda

Load and explore dataset: show shape, numeric columns, distributions, and correlation heatmap

```
[3]: raw_df = pd.read_csv(DATA_PATH)
print(f"Shape of dataset: {raw_df.shape}")
display(raw_df.head())

numeric_df = raw_df.select_dtypes(include=["number"])

raw_df.info()
display(raw_df.count(numeric_only=True).to_frame(name="count"))
print("Dataset shape:", raw_df.shape)

for col in numeric_df.columns:
    series = numeric_df[col]
    counts = series.value_counts(dropna=False).sort_index()
    x = range(len(counts))
    labels = counts.index.astype(str)
    plt.figure(figsize=(8, 3))
    plt.bar(x, counts.values, color="steelblue", alpha=0.9, width=0.8)
    plt.title(f"Count per value - {col}", fontsize=12, pad=8)
    plt.xlabel("Value")
    plt.ylabel("Count")
    if len(labels) > 40:
        step = max(1, len(labels) // 40)
        plt.xticks(list(x)[::step], labels[::step], rotation=45, ha="right")
    else:
        plt.xticks(x, labels, rotation=45, ha="right")
    plt.grid(True, linestyle="--", axis="y", alpha=0.4)
    plt.tight_layout()
    plt.show()

corr = numeric_df.corr()
plt.figure(figsize=(10, 10))
sns.heatmap(
    corr,
    annot=True,
    cmap="coolwarm",
    fmt=".1f",
    square=True,
```

```

    linewidths=1,
    cbar_kws={"shrink": 0.7, "label": "Correlation"}
)
plt.title("Correlation Heatmap of Numeric Attributes", fontsize=14, pad=12)
plt.tight_layout()
plt.show()

```

Shape of dataset: (1048575, 21)

	USMER	MEDICAL_UNIT	SEX	PATIENT_TYPE	DATE_DIED	INTUBED	PNEUMONIA	\	
0	2		1	1	1	03/05/2020	97	1	
1	2		1	2	1	03/06/2020	97	1	
2	2		1	2	2	09/06/2020	1	2	
3	2		1	1	1	12/06/2020	97	2	
4	2		1	2	1	21/06/2020	97	2	
	AGE	PREGNANT	DIABETES	...	ASTHMA	INMSUPR	HIPERTENSION	OTHER_DISEASE	\
0	65	2	2	...	2	2	1	2	
1	72	97	2	...	2	2	1	2	
2	55	97	1	...	2	2	2	2	
3	53	2	2	...	2	2	2	2	
4	68	97	1	...	2	2	1	2	
	CARDIOVASCULAR	OBESITY	RENAL_CHRONIC		TOBACCO	CLASIFFICATION_FINAL	ICU		
0		2	2		2	2	3	97	
1		2	1		1	2	5	97	
2		2	2		2	2	3	2	
3		2	2		2	2	7	97	
4		2	2		2	2	3	97	

[5 rows x 21 columns]

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1048575 entries, 0 to 1048574
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   USMER            1048575 non-null   int64  
 1   MEDICAL_UNIT     1048575 non-null   int64  
 2   SEX               1048575 non-null   int64  
 3   PATIENT_TYPE     1048575 non-null   int64  
 4   DATE_DIED        1048575 non-null   object 
 5   INTUBED           1048575 non-null   int64  
 6   PNEUMONIA         1048575 non-null   int64  
 7   AGE               1048575 non-null   int64  
 8   PREGNANT          1048575 non-null   int64  
 9   DIABETES          1048575 non-null   int64  
 10  COPD              1048575 non-null   int64  
 11  ASTHMA            1048575 non-null   int64

```

```


12 INMSUPR           1048575 non-null  int64
13 HIPERTENSION      1048575 non-null  int64
14 OTHER_DISEASE     1048575 non-null  int64
15 CARDIOVASCULAR    1048575 non-null  int64
16 OBESITY            1048575 non-null  int64
17 RENAL_CHRONIC     1048575 non-null  int64
18 TOBACCO            1048575 non-null  int64
19 CLASIFICATION_FINAL 1048575 non-null  int64
20 ICU                1048575 non-null  int64
dtypes: int64(20), object(1)
memory usage: 168.0+ MB




```

	count
USMER	1048575
MEDICAL_UNIT	1048575
SEX	1048575
PATIENT_TYPE	1048575
INTUBED	1048575
PNEUMONIA	1048575
AGE	1048575
PREGNANT	1048575
DIABETES	1048575
COPD	1048575
ASTHMA	1048575
INMSUPR	1048575
HIPERTENSION	1048575
OTHER_DISEASE	1048575
CARDIOVASCULAR	1048575
OBESITY	1048575
RENAL_CHRONIC	1048575
TOBACCO	1048575
CLASIFICATION_FINAL	1048575
ICU	1048575


Dataset shape: (1048575, 21)

Count per value — USMER







Count per value — INTUBED









Count per value — PNEUMONIA








Count per value — AGE















1.1 Exploratory analysis

[4]: `display(raw_df.describe(include='all').transpose())`

	count	unique	top	freq	mean	\
USMER	1048575.0	NaN	NaN	NaN	1.632194	
MEDICAL_UNIT	1048575.0	NaN	NaN	NaN	8.980565	
SEX	1048575.0	NaN	NaN	NaN	1.499259	
PATIENT_TYPE	1048575.0	NaN	NaN	NaN	1.190765	
DATE_DIED	1048575	401	9999-99-99	971633		NaN
INTUBED	1048575.0	NaN	NaN	NaN	79.522875	
PNEUMONIA	1048575.0	NaN	NaN	NaN	3.346831	
AGE	1048575.0	NaN	NaN	NaN	41.794102	
PREGNANT	1048575.0	NaN	NaN	NaN	49.765585	
DIABETES	1048575.0	NaN	NaN	NaN	2.186404	
COPD	1048575.0	NaN	NaN	NaN	2.260569	
ASTHMA	1048575.0	NaN	NaN	NaN	2.242626	
INMSUPR	1048575.0	NaN	NaN	NaN	2.298132	
HIPERTENSION	1048575.0	NaN	NaN	NaN	2.128989	
OTHER_DISEASE	1048575.0	NaN	NaN	NaN	2.435143	
CARDIOVASCULAR	1048575.0	NaN	NaN	NaN	2.26181	
OBESITY	1048575.0	NaN	NaN	NaN	2.125176	
RENAL_CHRONIC	1048575.0	NaN	NaN	NaN	2.25718	
TOBACCO	1048575.0	NaN	NaN	NaN	2.214333	
CLASIFICATION_FINAL	1048575.0	NaN	NaN	NaN	5.305653	
ICU	1048575.0	NaN	NaN	NaN	79.553974	


	std	min	25%	50%	75%	max
USMER	0.482208	1.0	1.0	2.0	2.0	2.0
MEDICAL_UNIT	3.723278	1.0	4.0	12.0	12.0	13.0
SEX	0.5	1.0	1.0	1.0	2.0	2.0
PATIENT_TYPE	0.392904	1.0	1.0	1.0	1.0	2.0
DATE_DIED	NaN	NaN	NaN	NaN	NaN	NaN
INTUBED	36.868886	1.0	97.0	97.0	97.0	99.0
PNEUMONIA	11.912881	1.0	2.0	2.0	2.0	99.0
AGE	16.907389	0.0	30.0	40.0	53.0	121.0
PREGNANT	47.510733	1.0	2.0	97.0	97.0	98.0
DIABETES	5.424242	1.0	2.0	2.0	2.0	98.0
COPD	5.132258	1.0	2.0	2.0	2.0	98.0
ASTHMA	5.114089	1.0	2.0	2.0	2.0	98.0
INMSUPR	5.462843	1.0	2.0	2.0	2.0	98.0
HIPERTENSION	5.236397	1.0	2.0	2.0	2.0	98.0
OTHER_DISEASE	6.646676	1.0	2.0	2.0	2.0	98.0
CARDIOVASCULAR	5.19485	1.0	2.0	2.0	2.0	98.0
OBESITY	5.175445	1.0	2.0	2.0	2.0	98.0
RENAL_CHRONIC	5.135354	1.0	2.0	2.0	2.0	98.0
TOBACCO	5.323097	1.0	2.0	2.0	2.0	98.0
CLASIFICATION_FINAL	1.881165	1.0	3.0	6.0	7.0	7.0

```
ICU           36.823073  1.0  97.0  97.0  97.0  99.0
```


Create binary target (1 = COVID, 0 = No COVID) and visualize target and age distributions

```
[5]: target_series = raw_df['CLASIFICATION_FINAL'].isin([1, 2, 3]).astype(int)
fig, axes = plt.subplots(1, 2, figsize=(10, 4))
target_series.value_counts(normalize=True).mul(100).reindex([0, 1]).plot(
    kind='bar', ax=axes[0], color=['steelblue', 'tomato'])
)
axes[0].set_title('Target distribution (%)')
axes[0].set_xlabel('HAS_COVID')
axes[0].set_ylabel('Percent')

age_sample = raw_df['AGE'].clip(0, 100)
axes[1].hist(age_sample, bins=30, color='slategray', alpha=0.9)
axes[1].set_title('Age distribution (trimmed 0-100)')
axes[1].set_xlabel('Age')
axes[1].set_ylabel('Count')
plt.tight_layout()
plt.show()
```



```
[6]: numeric_cols = raw_df.select_dtypes(include=['number']).columns
sample_for_corr = raw_df[numeric_cols].sample(n=min(2000, len(raw_df)), random_state=42)
plt.figure(figsize=(10, 8))
sns.heatmap(sample_for_corr.corr(), cmap='coolwarm', linewidths=0.1, center=0)
plt.title('Correlation heatmap (raw numeric sample)')
plt.tight_layout()
plt.show()
```



```
[7]: demo_df = raw_df.copy()
print('Sample before preprocessing:')
display(demo_df.iloc[:30])

demo_df.columns = demo_df.columns.str.strip()
demo_df.rename(columns={'DATE_DIED': 'DEAD'}, inplace=True)
med_unit_original = demo_df['MEDICAL_UNIT'].copy()
classif_original = demo_df['CLASIFICATION_FINAL'].copy()

demo_df['DEAD'] = np.where(demo_df['DEAD'] == '9999-99-99', 2, 1)
demo_df.loc[demo_df['SEX'] == 1, 'PREGNANT'] = 0
binary_cols = [
    'USMER', 'SEX', 'INTUBED', 'PNEUMONIA', 'PREGNANT', 'DIABETES', 'COPD', 'ASTHMA', 'INMSUPR', 'HYPERTENSION',
    'OTHER_DISEASE', 'CARDIOVASCULAR', 'OBESITY', 'RENAL_CHRONIC', 'TOBACCO'
]
demo_df[binary_cols] = demo_df[binary_cols].replace(2, 0)
demo_df[binary_cols] = demo_df[binary_cols].replace(97, 0.5)
demo_df[binary_cols] = demo_df[binary_cols].replace(98, 0.5)
demo_df['MEDICAL_UNIT'] = med_unit_original
demo_df['CLASIFICATION_FINAL'] = classif_original
```

```

demo_df.rename(columns={'SEX': 'IS_MALE'}, inplace=True)
demo_df['CLASIFICATION_FINAL'] = demo_df['CLASIFICATION_FINAL'].
    ↪replace([1,2,3], 1)
demo_df['CLASIFICATION_FINAL'] = demo_df['CLASIFICATION_FINAL'].
    ↪replace([4,5,6,7], 0)
demo_df.rename(columns={'CLASIFICATION_FINAL': 'HAS_COVID'}, inplace=True)
print('Sample after key preprocessing steps:')
display(demo_df.iloc[:30])
# Display purpose only

```

Sample before preprocessing:

	USMER	MEDICAL_UNIT	SEX	PATIENT_TYPE	DATE_DIED	INTUBED	PNEUMONIA	\
0	2		1	1	1 03/05/2020	97	1	
1	2		1	2	1 03/06/2020	97	1	
2	2		1	2	2 09/06/2020	1	2	
3	2		1	1	1 12/06/2020	97	2	
4	2		1	2	1 21/06/2020	97	2	
5	2		1	1	2 9999-99-99	2	1	
6	2		1	1	1 9999-99-99	97	2	
7	2		1	1	1 9999-99-99	97	1	
8	2		1	1	2 9999-99-99	2	2	
9	2		1	1	2 9999-99-99	2	2	
10	2		1	1	1 9999-99-99	97	2	
11	2		1	2	2 9999-99-99	2	2	
12	2		1	2	2 9999-99-99	2	2	
13	2		1	2	1 9999-99-99	97	2	
14	2		1	1	1 9999-99-99	97	2	
15	2		1	1	1 9999-99-99	97	2	
16	2		1	1	2 9999-99-99	2	1	
17	2		1	2	1 9999-99-99	97	2	
18	2		1	2	1 9999-99-99	97	2	
19	2		1	1	1 9999-99-99	97	2	
20	2		1	2	2 9999-99-99	2	1	
21	2		1	2	1 9999-99-99	97	2	
22	2		1	2	1 9999-99-99	97	2	
23	2		1	1	1 9999-99-99	97	2	
24	2		1	1	1 9999-99-99	97	2	
25	2		1	2	1 9999-99-99	97	2	
26	2		1	2	1 9999-99-99	97	2	
27	2		1	2	1 9999-99-99	97	2	
28	2		1	2	1 9999-99-99	97	2	
29	2		1	2	1 9999-99-99	97	2	
	AGE	PREGNANT	DIABETES	...	ASTHMA	INMSUPR	HIPERTENSION	\
0	65	2	2	...	2	2	1	
1	72	97	2	...	2	2	1	
2	55	97	1	...	2	2	2	

3	53	2	2	...	2	2	2
4	68	97	1	...	2	2	1
5	40	2	2	...	2	2	2
6	64	2	2	...	2	2	2
7	64	2	1	...	2	1	1
8	37	2	1	...	2	2	1
9	25	2	2	...	2	2	2
10	38	2	2	...	2	2	2
11	24	97	2	...	2	2	2
12	30	97	2	...	2	2	2
13	55	97	2	...	2	2	2
14	48	2	1	...	2	2	2
15	23	2	2	...	2	2	2
16	80	2	2	...	2	2	1
17	61	97	2	...	2	2	2
18	54	97	2	...	2	2	2
19	64	2	2	...	2	2	2
20	59	97	1	...	2	2	2
21	30	97	2	...	2	2	2
22	45	97	2	...	2	2	2
23	26	2	2	...	2	2	2
24	38	2	2	...	2	2	2
25	24	97	2	...	2	2	2
26	32	97	2	...	2	2	2
27	49	97	2	...	2	2	2
28	39	97	2	...	2	2	2
29	27	97	2	...	2	2	2

	OTHER_DISEASE	CARDIOVASCULAR	OBESITY	RENAL_CHRONIC	TOBACCO	\
0	2	2	2	2	2	2
1	2	2	1	1	1	2
2	2	2	2	2	2	2
3	2	2	2	2	2	2
4	2	2	2	2	2	2
5	2	2	2	2	2	2
6	2	2	2	2	2	2
7	2	2	2	1	1	2
8	2	2	1	2	2	2
9	2	2	2	2	2	2
10	2	2	2	2	2	2
11	2	2	2	2	2	2
12	2	2	2	2	2	2
13	2	2	2	2	2	2
14	2	2	2	2	2	2
15	2	2	2	2	2	2
16	2	2	2	2	2	2
17	2	2	2	2	2	2
18	2	2	2	2	2	2

19	2	2	2	2	2
20	2	2	2	2	1
21	2	2	2	2	2
22	2	2	2	2	2
23	2	2	2	2	2
24	2	2	2	1	2
25	2	2	2	2	2
26	2	2	2	2	2
27	2	2	2	2	2
28	2	2	2	2	1
29	2	2	2	2	2

	CLASIFFICATION_FINAL	ICU
0	3	97
1	5	97
2	3	2
3	7	97
4	3	97
5	3	2
6	3	97
7	3	97
8	3	2
9	3	2
10	3	97
11	3	2
12	3	2
13	3	97
14	3	97
15	3	97
16	3	1
17	3	97
18	3	97
19	3	97
20	3	1
21	3	97
22	3	97
23	3	97
24	3	97
25	3	97
26	3	97
27	3	97
28	3	97
29	3	97

[30 rows x 21 columns]

Sample after key preprocessing steps:

USMER	MEDICAL_UNIT	IS_MALE	PATIENT_TYPE	DEAD	INTUBED	PNEUMONIA	AGE	\
-------	--------------	---------	--------------	------	---------	-----------	-----	---

0	0	1	1	1	1	0.5	1	65
1	0	1	0	1	1	0.5	1	72
2	0	1	0	2	1	1.0	0	55
3	0	1	1	1	1	0.5	0	53
4	0	1	0	1	1	0.5	0	68
5	0	1	1	2	0	0.0	1	40
6	0	1	1	1	0	0.5	0	64
7	0	1	1	1	0	0.5	1	64
8	0	1	1	2	0	0.0	0	37
9	0	1	1	2	0	0.0	0	25
10	0	1	1	1	0	0.5	0	38
11	0	1	0	2	0	0.0	0	24
12	0	1	0	2	0	0.0	0	30
13	0	1	0	1	0	0.5	0	55
14	0	1	1	1	0	0.5	0	48
15	0	1	1	1	0	0.5	0	23
16	0	1	1	2	0	0.0	1	80
17	0	1	0	1	0	0.5	0	61
18	0	1	0	1	0	0.5	0	54
19	0	1	1	1	0	0.5	0	64
20	0	1	0	2	0	0.0	1	59
21	0	1	0	1	0	0.5	0	30
22	0	1	0	1	0	0.5	0	45
23	0	1	1	1	0	0.5	0	26
24	0	1	1	1	0	0.5	0	38
25	0	1	0	1	0	0.5	0	24
26	0	1	0	1	0	0.5	0	32
27	0	1	0	1	0	0.5	0	49
28	0	1	0	1	0	0.5	0	39
29	0	1	0	1	0	0.5	0	27

PREGNANT	DIABETES	...	ASTHMA	INMSUPR	HIPERTENSION	OTHER_DISEASE	\
0	0.0	0.0	...	0.0	0.0	1.0	0.0
1	0.5	0.0	...	0.0	0.0	1.0	0.0
2	0.5	1.0	...	0.0	0.0	0.0	0.0
3	0.0	0.0	...	0.0	0.0	0.0	0.0
4	0.5	1.0	...	0.0	0.0	1.0	0.0
5	0.0	0.0	...	0.0	0.0	0.0	0.0
6	0.0	0.0	...	0.0	0.0	0.0	0.0
7	0.0	1.0	...	0.0	1.0	1.0	0.0
8	0.0	1.0	...	0.0	0.0	1.0	0.0
9	0.0	0.0	...	0.0	0.0	0.0	0.0
10	0.0	0.0	...	0.0	0.0	0.0	0.0
11	0.5	0.0	...	0.0	0.0	0.0	0.0
12	0.5	0.0	...	0.0	0.0	0.0	0.0
13	0.5	0.0	...	0.0	0.0	0.0	0.0
14	0.0	1.0	...	0.0	0.0	0.0	0.0
15	0.0	0.0	...	0.0	0.0	0.0	0.0

16	0.0	0.0	...	0.0	0.0	1.0	0.0
17	0.5	0.0	...	0.0	0.0	0.0	0.0
18	0.5	0.0	...	0.0	0.0	0.0	0.0
19	0.0	0.0	...	0.0	0.0	0.0	0.0
20	0.5	1.0	...	0.0	0.0	0.0	0.0
21	0.5	0.0	...	0.0	0.0	0.0	0.0
22	0.5	0.0	...	0.0	0.0	0.0	0.0
23	0.0	0.0	...	0.0	0.0	0.0	0.0
24	0.0	0.0	...	0.0	0.0	0.0	0.0
25	0.5	0.0	...	0.0	0.0	0.0	0.0
26	0.5	0.0	...	0.0	0.0	0.0	0.0
27	0.5	0.0	...	0.0	0.0	0.0	0.0
28	0.5	0.0	...	0.0	0.0	0.0	0.0
29	0.5	0.0	...	0.0	0.0	0.0	0.0

	CARDIOVASCULAR	OBESITY	RENAL_CHRONIC	TOBACCO	HAS_COVID	ICU
0	0.0	0.0	0.0	0.0	1	0.5
1	0.0	1.0	1.0	0.0	0	0.5
2	0.0	0.0	0.0	0.0	1	0.0
3	0.0	0.0	0.0	0.0	0	0.5
4	0.0	0.0	0.0	0.0	1	0.5
5	0.0	0.0	0.0	0.0	1	0.0
6	0.0	0.0	0.0	0.0	1	0.5
7	0.0	0.0	1.0	0.0	1	0.5
8	0.0	1.0	0.0	0.0	1	0.0
9	0.0	0.0	0.0	0.0	1	0.0
10	0.0	0.0	0.0	0.0	1	0.5
11	0.0	0.0	0.0	0.0	1	0.0
12	0.0	0.0	0.0	0.0	1	0.0
13	0.0	0.0	0.0	0.0	1	0.5
14	0.0	0.0	0.0	0.0	1	0.5
15	0.0	0.0	0.0	0.0	1	0.5
16	0.0	0.0	0.0	0.0	1	1.0
17	0.0	0.0	0.0	0.0	1	0.5
18	0.0	0.0	0.0	0.0	1	0.5
19	0.0	0.0	0.0	0.0	1	0.5
20	0.0	0.0	0.0	1.0	1	1.0
21	0.0	0.0	0.0	0.0	1	0.5
22	0.0	0.0	0.0	0.0	1	0.5
23	0.0	0.0	0.0	0.0	1	0.5
24	0.0	0.0	1.0	0.0	1	0.5
25	0.0	0.0	0.0	0.0	1	0.5
26	0.0	0.0	0.0	0.0	1	0.5
27	0.0	0.0	0.0	0.0	1	0.5
28	0.0	0.0	0.0	1.0	1	0.5
29	0.0	0.0	0.0	0.0	1	0.5

[30 rows x 21 columns]

1.2 Preprocessing

Our preprocessing function simply prepares the dataframe we have, for model. We work with only valid values, retype non-numeric values to numeric, normalize data by leaving 1 (true) be 1, but 2 (false) to be 0 and unknown values with np.nan. We ensure that male patients cannot have pregnant value set to 1. # Most importantly we replace values in final classification from 1-6 to be only 1 or 2 redefining it to binary classification. (it now represents if the patient is sick, or is not sick). And we also normalize the age of patients

```
[8]: # Preprocessing function that we actually use
def preprocess_data(frame: pd.DataFrame) -> pd.DataFrame:
    df = frame.copy() # work on a copy to avoid modifying the original
    ↵dataframe

    # Clean column names and rename gender column for clarity
    df.columns = df.columns.str.strip()
    df.rename(columns={"SEX": "IS_MALE"}, inplace=True)

    # Keep only rows with valid COVID classification values (1-6)
    valid = df["CLASIFICATION_FINAL"].isin([1, 2, 3, 4, 5, 6])
    df = df.loc[valid].copy()

    # Create binary column DEAD (1 = dead, 0 = alive)
    df["DEAD"] = np.where(df["DATE_DIED"] == "9999-99-99", 0, 1)

    # List of binary features (0/1 type)
    binary_cols = [
        "USMER", "IS_MALE", "INTUBED", "PNEUMONIA", "PREGNANT", "DIABETES",
        ↵"COPD",
        "ASTHMA", "INMSUPR", "HIPERTENSION", "OTHER_DISEASE", "CARDIOVASCULAR",
        "OBESITY", "RENAL_CHRONIC", "TOBACCO", "ICU", "DEAD",
    ]
    ↵]

    # Replace invalid codes (97-99 etc.) with NaN to mark missing data
    replace_map = {1: 1, 2: 0, 3: np.nan, 97: np.nan, 98: np.nan, 99: np.nan}
    df[binary_cols] = df[binary_cols].replace(replace_map)

    # Ensure males are never marked as pregnant
    df.loc[df["IS_MALE"] == 1, "PREGNANT"] = 0

    # Create hospitalization flag: 1 = hospitalized, 0 = outpatient
    df["IS_HOSPITALIZED"] = df["PATIENT_TYPE"].map({1: 0, 2: 1}).fillna(0)
    df.drop(columns=["PATIENT_TYPE"], inplace=True)

    # Convert COVID classification (1-3 positive, 4-6 negative) into binary
    ↵label
    df["HAS_COVID"] = df["CLASIFICATION_FINAL"].replace({1: 1, 2: 1, 3: 1, 4:
    ↵0, 5: 0, 6: 0})
```

```

df.drop(columns=["CLASIFICATION_FINAL", "DATE_DIED"], inplace=True)

# Fill missing binary values with median (usually 0 or 1)
df[binary_cols] = df[binary_cols].fillna(df[binary_cols].median())

# Define chronic condition columns
risk_cols = [
    "DIABETES", "COPD", "ASTHMA", "INMSUPR", "HIPERTENSION",
    "OTHER_DISEASE", "CARDIOVASCULAR", "OBESITY", "RENAL_CHRONIC",
]
]

# Create aggregated health indicators
df["CHRONIC_COUNT"] = df[risk_cols].sum(axis=1) # number of chronic diseases
df["MULTI_MORBID"] = (df["CHRONIC_COUNT"] >= 2).astype(int) # has 2+ chronic diseases
df["RISK_OBESE_SMOKER"] = ((df["OBESITY"] == 1) | (df["TOBACCO"] == 1)).astype(int) # obese or smoker

# Bucketize age into groups (0-30, 31-45, 46-60, 61-75, 76+)
df["AGE_BUCKET"] = (
    pd.cut(df["AGE"], bins=[0, 30, 45, 60, 75, 120], labels=False, include_lowest=True)
    .fillna(0)
    .astype(int)
)

# Normalize age using standard scaling (mean=0, std=1)
scaler = StandardScaler()
df["AGE"] = scaler.fit_transform(df[["AGE"]])

# One-hot encode medical unit column (convert category → binary columns)
df = pd.get_dummies(df, columns=["MEDICAL_UNIT"], prefix="MED_UNIT", drop_first=True)

# Drop any remaining missing values and reset index
df = df.dropna()
df.reset_index(drop=True, inplace=True)

# Return cleaned and fully preprocessed dataset
return df

```

```
[9]: model_df = preprocess_data(raw_df)
print(f"Processed shape: {model_df.shape}")
display(model_df.head())
feature_cols = [col for col in model_df.columns if col != "HAS_COVID"]
print(f"Feature count: {len(feature_cols)}")
```

Processed shape: (549325, 36)

```
    USMER  IS_MALE  INTUBED  PNEUMONIA      AGE  PREGNANT  DIABETES  COPD  \
0      0         1        0.0       1.0  1.262042      0.0        0.0     0.0
1      0         0        0.0       1.0  1.676426      0.0        0.0     0.0
2      0         0        1.0       0.0  0.670064      0.0        1.0     0.0
3      0         0        0.0       0.0  1.439635      0.0        1.0     0.0
4      0         1        0.0       1.0 -0.217902      0.0        0.0     0.0


  ASTHMA  INMSUPR  ...  MED_UNIT_4  MED_UNIT_5  MED_UNIT_6  MED_UNIT_7  \
0     0.0     0.0  ...     False     False     False     False
1     0.0     0.0  ...     False     False     False     False
2     0.0     0.0  ...     False     False     False     False
3     0.0     0.0  ...     False     False     False     False
4     0.0     0.0  ...     False     False     False     False

  MED_UNIT_8  MED_UNIT_9  MED_UNIT_10  MED_UNIT_11  MED_UNIT_12  MED_UNIT_13
0     False     False     False     False     False     False     False
1     False     False     False     False     False     False     False
2     False     False     False     False     False     False     False
3     False     False     False     False     False     False     False
4     False     False     False     False     False     False     False
```

[5 rows x 36 columns]

Feature count: 35

```
[10]: processed_numeric = model_df.select_dtypes(include=['number'])
corr_processed = processed_numeric.corr()
plt.figure(figsize=(10, 10))
sns.heatmap(
    corr_processed,
    annot=True,
    cmap='coolwarm',
    fmt='.1f',
    square=True,
    linewidths=1,
    cbar_kws={'shrink': 0.7, 'label': 'Correlation'}
)
plt.title('Correlation Heatmap after Preprocessing', fontsize=14, pad=12)
plt.tight_layout()
plt.show()
```



1.3 Data split

Next we split data for training (70%) for validation (10%) and for testing (20%) so we can accurately train our NN

```
[11]: # Split dataset into train (70%), validation (10%), and test (20%) sets
train_df, temp_df = train_test_split(
    model_df,
    test_size=0.3,                                     # 30% goes to temp (val + test)
    stratify=model_df["HAS_COVID"],                   # keep same class ratio across
    random_state=42,                                   # ensure reproducibility
    shuffle=True,
)

# Split temp into validation (1/3 of temp = 10%) and test (2/3 of temp = 20%)
```

```

val_df, test_df = train_test_split(
    temp_df,
    test_size=2 / 3,
    stratify=temp_df["HAS_COVID"],
    random_state=42,
    shuffle=True,
)

# Reset indices for clean dataframes
train_df = train_df.reset_index(drop=True)
val_df = val_df.reset_index(drop=True)
test_df = test_df.reset_index(drop=True)

# Helper function to print class distribution info
def describe_split(name, frame):
    share = frame["HAS_COVID"].mean() * 100
    print(f"{name}:<10s} n={len(frame):,} | positive={share:.2f}% | negative={100 - share:.2f}%")

# Display dataset split summary
print("Dataset splits:")
describe_split("Train", train_df)
describe_split("Val", val_df)
describe_split("Test", test_df)

```

Dataset splits:

Train	n=384,527	positive=71.36%	negative=28.64%
Val	n=54,932	positive=71.36%	negative=28.64%
Test	n=109,866	positive=71.36%	negative=28.64%

1.4 Torch datasets

From now on we will work with PyTorch, so we need to convert our DataFrames to TensorDatasets

```
[12]: # Convert DataFrame to PyTorch TensorDataset (features + labels)
def frame_to_dataset(frame):
    features = frame[feature_cols].to_numpy(dtype=np.float32)      # input features
    labels = frame["HAS_COVID"].to_numpy(dtype=np.float32)          # target labels
    return TensorDataset(torch.from_numpy(features), torch.from_numpy(labels))

# Create datasets for training, validation, and testing
train_dataset = frame_to_dataset(train_df)
val_dataset = frame_to_dataset(val_df)
test_dataset = frame_to_dataset(test_df)

# Compute class weights to handle imbalance (more weight for minority class)
class_counts = train_df["HAS_COVID"].value_counts().to_dict()
```

```

weights = train_df["HAS_COVID"].map(lambda value: 1.0 / class_counts[value]).  

    ↪to_numpy(dtype=np.float32)

# Create a sampler for balanced batches during training
sampler = WeightedRandomSampler(weights=torch.from_numpy(weights),  

    ↪num_samples=len(weights), replacement=True)

# Build data loaders for efficient batching
train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, sampler=sampler)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE)

# Calculate positive class weight for loss function (to handle imbalance)
positive = float((train_df["HAS_COVID"] == 1).sum())
negative = float((train_df["HAS_COVID"] == 0).sum())
pos_weight_value = negative / max(positive, 1.0)

# Define binary classification loss with class weight
criterion = nn.BCEWithLogitsLoss(pos_weight=torch.tensor([pos_weight_value]),  

    ↪dtype=torch.float32, device=device)

# Inspect one batch from the training loader
sample_batch = next(iter(train_loader))
print(f"Sample batch → X: {sample_batch[0].shape}, y: {sample_batch[1].shape},  

    ↪pos_weight={pos_weight_value:.2f}")

```

Sample batch → X: torch.Size([1024, 35]), y: torch.Size([1024]), pos_weight=0.40

1.5 Model and helpers

Using PyTorch integrated linear function, activation function and other cool things, we create our model specialised for binary classification, predicting if the patient has Covid or not. # We also calculate metrics like f1 score to determine how good is our NN trained.

```
[13]: # Simple feedforward MLP for binary COVID classification
class CovidMLP(nn.Module):
    def __init__(self, input_dim, hidden_layers, dropout):
        super().__init__()
        layers = []
        prev = input_dim # start with the input feature size

        # build hidden layers based on config
        for hidden in hidden_layers:
            layers.append(nn.Linear(prev, hidden)) # dense layer
            layers.append(nn.BatchNorm1d(hidden)) # helps stabilize
    ↪training
            layers.append(nn.ReLU()) # non-linearity
            layers.append(nn.Dropout(dropout)) # regularization
```

```

        prev = hidden                                # next layer input =_
        ↵this layer output

        # final output layer (1 neuron for binary classification)
        layers.append(nn.Linear(prev, 1))
        self.net = nn.Sequential(*layers)

    def forward(self, x):
        # forward pass through the network
        return self.net(x).squeeze(1) # remove extra dimension (batch, 1) →_
        ↵(batch,)

# Calculate common metrics from probabilities
def metrics_from_probs(probs, labels, threshold=0.5):
    preds = (probs >= threshold).astype(int) # turn probs into 0/1 predictions

    metrics = {
        "accuracy": accuracy_score(labels, preds),
        "precision": precision_score(labels, preds, zero_division=0),
        "recall": recall_score(labels, preds, zero_division=0),
        "f1": f1_score(labels, preds, zero_division=0),
        "pr_auc": average_precision_score(labels, probs), # area under_
        ↵precision-recall curve
    }

    # how much recall we get when precision is at least 0.7
    prec_curve, rec_curve, _ = precision_recall_curve(labels, probs)
    mask = prec_curve >= 0.7
    metrics["recall_at_70_precision"] = float(rec_curve[mask].max()) if mask.
    ↵any() else 0.0

    return metrics

# Run evaluation on val/test data and collect metrics
def evaluate_model(model, loader):
    model.eval() # disable dropout, switch to eval mode
    losses, probs, labels = [], [], []

    with torch.no_grad(): # no gradients needed for evaluation
        for xb, yb in loader:
            xb = xb.to(device)
            yb = yb.to(device)

            logits = model(xb)           # forward pass
            loss = criterion(logits, yb) # compute loss

```

```

        losses.append(loss.item())

        # convert logits → probabilities and collect labels
        probs.append(torch.sigmoid(logits).cpu())
        labels.append(yb.cpu())

        # merge all batches into full arrays
        prob_array = torch.cat(probs).numpy()
        label_array = torch.cat(labels).numpy()

        # compute metrics
        metrics = metrics_from_probs(prob_array, label_array)
        metrics["roc_auc"] = roc_auc_score(label_array, prob_array)
        metrics["loss"] = float(np.mean(losses)) # average val loss

    return metrics, prob_array, label_array

# Find the threshold that gives the best F1 score
def find_best_threshold(probs, labels):
    best_thr = 0.5
    best_score = 0.0

    # try thresholds from 0.1 to 0.9
    for thr in np.linspace(0.1, 0.9, 17):
        score = f1_score(labels, (probs >= thr).astype(int), zero_division=0)
        if score > best_score:
            best_score = score
            best_thr = float(thr)

    return best_thr, best_score

```

Our model training function builds the model from the config, initializes the optimizer, scheduler, and logs everything using Wandb. # Then it runs set amount of epochs from config, training the model, changing the learning rate and logging more data, and saving the best result for later

```
[14]: def train_model(config):
    # Build model from config and move to device
    model = CovidMLP(len(feature_cols), config["hidden_layers"], ↴
    config["dropout"]).to(device)

    # Optimizer: AdamW (Adam + decoupled weight decay for better generalization)
    optimizer = torch.optim.AdamW(
        model.parameters(),
        lr=config["lr"],
        weight_decay=config["weight_decay"]
    )
```

```

# LR scheduler: reduce LR when validation loss plateaus (mode='min')
# factor=0.5 → halve LR; patience=1 → wait 1 epoch without improvement
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer, mode="min", factor=0.5, patience=1
)

# Initialize Weights & Biases run for experiment tracking
# - project: group runs
# - config: hyperparams logged as run config
# - reinit=False: reuse the same process/run context
wandb_run = wandb.init(project="covid-19", config=config, reinit=False)
wandb_run.watch(model, log="all") # log gradients/weights for debugging

best_state = None # best model weights (by val F1)
best_snapshot = None # (metrics, probs, labels) for best epoch
best_score = -np.inf # track best F1
history = [] # per-epoch metrics for plotting

# ---- Training loop ----
for epoch in range(1, config["epochs"] + 1):
    model.train() # enable dropout/batchnorm updates
    epoch_losses = []

    # Iterate mini-batches
    for xb, yb in train_loader:
        xb = xb.to(device)
        yb = yb.to(device)

        optimizer.zero_grad() # reset accumulated grads
        logits = model(xb) # forward pass → logits
        loss = criterion(logits, yb) # BCEWithLogitsLoss (with
        ↵pos_weight)
        loss.backward() # backpropagate gradients
        optimizer.step() # update weights

    epoch_losses.append(loss.item())

# ---- Validation ----
val_metrics, val_probs, val_labels = evaluate_model(model, val_loader)

# Step LR scheduler with validation loss (expects a scalar)
scheduler.step(val_metrics["loss"])

# Save epoch summary for later analysis/plots
history.append({
    "epoch": epoch,

```

```

        "train_loss": float(np.mean(epoch_losses)),
        **{k: v for k, v in val_metrics.items()},
    })

# Console log
print(
    f"[config['name']:<10s] epoch {epoch:02d} "
    f"train_loss={history[-1]['train_loss']:.3f} "
    f"val_loss={val_metrics['loss']:.3f} "
    f"val_f1={val_metrics['f1']:.3f}"
)

# Log key metrics to W&B for dashboards/curves
wandb_run.log({
    "epoch": epoch,
    "train_loss": float(np.mean(epoch_losses)),
    "val_loss": val_metrics["loss"],
    "val_f1": val_metrics["f1"],
    "learning_rate": optimizer.param_groups[0]["lr"],
})

# Track best model by validation F1
if val_metrics["f1"] > best_score:
    best_score = val_metrics["f1"]
    best_state = deepcopy(model.state_dict()) # keep ↴weights
    best_snapshot = (val_metrics.copy(), val_probs.copy(), # keep ↴artifacts
                     val_labels.copy())

# Restore best weights before returning
model.load_state_dict(best_state)

# Close W&B run cleanly
wandb_run.finish()

# Return trained (best) model, full training history, and best-epoch ↴snapshot
return model, history, best_snapshot

```

1.6 Experiments

We define a few experiments, with different names so we can track them later. Each of these experiments is a config for our model training function, after each training we save the results of the experiment so we can compare them later.

```
[15]: # Define a small grid of experiments (architectures + hyperparameters)
experiment_grid = [
    {"name": "baseline", "hidden_layers": [128, 64, 32], "dropout": 0.3, "lr": 3e-4, "weight_decay": 1e-4, "epochs": 10},
    {"name": "compact", "hidden_layers": [64, 32], "dropout": 0.25, "lr": 5e-4, "weight_decay": 5e-5, "epochs": 12},
    {"name": "wide", "hidden_layers": [256, 128, 64], "dropout": 0.4, "lr": 2e-4, "weight_decay": 1e-4, "epochs": 12},
]
experiment_results = [] # will store trained model snapshots + metrics for each config

# Run training for each config in the grid
for cfg in experiment_grid:
    model, history, snapshot = train_model(cfg) # train and return best-epoch snapshot
    metrics_at_05, val_probs, val_labels = snapshot # metrics at 0.5 threshold & raw probs/labels

    # Tune decision threshold to maximize F1 on validation set
    best_thr, _ = find_best_threshold(val_probs, val_labels)

    # Recompute metrics using the tuned threshold (keep ROC AUC & loss from snapshot)
    tuned_metrics = metrics_from_probs(val_probs, val_labels, threshold=best_thr)
    tuned_metrics["roc_auc"] = metrics_at_05["roc_auc"]
    tuned_metrics["loss"] = metrics_at_05["loss"]

    # Save experiment artifacts for later comparison / model selection
    experiment_results.append(
        {
            "name": cfg["name"], # experiment label
            "config": cfg, # full hyperparameter
            "state_dict": deepcopy(model.state_dict()), # best model weights
            "threshold": best_thr, # tuned decision threshold
            "val_metrics": tuned_metrics, # metrics at tuned threshold
            "history": history, # per-epoch training/val
        }
    )

# Build a compact comparison table across experiments
```

```

results_table = pd.DataFrame(
    [
        {
            "experiment": result["name"],
            "val_f1": result["val_metrics"]["f1"],
            "val_pr_auc": result["val_metrics"]["pr_auc"],
            "val_recall_at_70_precision": result["val_metrics"]["recall_at_70_precision"],
            "threshold": result["threshold"],
        }
        for result in experiment_results
    ]
).sort_values(by="val_f1", ascending=False) # rank by F1

print("Validation summary:")
display(results_table.reset_index(drop=True))

```

wandb: Currently logged in as: adamglog9919 (adamglog9919-my_org_ig) to https://api.wandb.ai. Use `wandb login --relogin` to force relogin
 wandb: WARNING Using a boolean value for 'reinit' is deprecated. Use 'return_previous' or 'finish_previous' instead.

```

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

baseline epoch 01 train_loss=0.447 val_loss=0.389 val_f1=0.367
baseline epoch 02 train_loss=0.395 val_loss=0.403 val_f1=0.266
baseline epoch 03 train_loss=0.390 val_loss=0.406 val_f1=0.261
baseline epoch 04 train_loss=0.389 val_loss=0.408 val_f1=0.254
baseline epoch 05 train_loss=0.388 val_loss=0.406 val_f1=0.255
baseline epoch 06 train_loss=0.387 val_loss=0.406 val_f1=0.260
baseline epoch 07 train_loss=0.388 val_loss=0.409 val_f1=0.243
baseline epoch 08 train_loss=0.387 val_loss=0.407 val_f1=0.251
baseline epoch 09 train_loss=0.387 val_loss=0.407 val_f1=0.247
baseline epoch 10 train_loss=0.387 val_loss=0.407 val_f1=0.247

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

```

```
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

compact    epoch 01 train_loss=0.403 val_loss=0.402 val_f1=0.253
compact    epoch 02 train_loss=0.389 val_loss=0.402 val_f1=0.243
compact    epoch 03 train_loss=0.387 val_loss=0.404 val_f1=0.245
compact    epoch 04 train_loss=0.386 val_loss=0.403 val_f1=0.253
compact    epoch 05 train_loss=0.385 val_loss=0.404 val_f1=0.249
compact    epoch 06 train_loss=0.385 val_loss=0.402 val_f1=0.247
compact    epoch 07 train_loss=0.385 val_loss=0.403 val_f1=0.243
compact    epoch 08 train_loss=0.385 val_loss=0.402 val_f1=0.245
compact    epoch 09 train_loss=0.385 val_loss=0.403 val_f1=0.253
compact    epoch 10 train_loss=0.384 val_loss=0.404 val_f1=0.248
compact    epoch 11 train_loss=0.383 val_loss=0.403 val_f1=0.246
compact    epoch 12 train_loss=0.384 val_loss=0.405 val_f1=0.240

<IPython.core.display.HTML object>

wide      epoch 01 train_loss=0.413 val_loss=0.399 val_f1=0.308
wide      epoch 02 train_loss=0.394 val_loss=0.401 val_f1=0.313
wide      epoch 03 train_loss=0.390 val_loss=0.403 val_f1=0.290
wide      epoch 04 train_loss=0.389 val_loss=0.403 val_f1=0.292
wide      epoch 05 train_loss=0.389 val_loss=0.401 val_f1=0.293
wide      epoch 06 train_loss=0.388 val_loss=0.402 val_f1=0.289
wide      epoch 07 train_loss=0.387 val_loss=0.403 val_f1=0.283
wide      epoch 08 train_loss=0.387 val_loss=0.403 val_f1=0.285
wide      epoch 09 train_loss=0.387 val_loss=0.403 val_f1=0.286
wide      epoch 10 train_loss=0.387 val_loss=0.403 val_f1=0.279
wide      epoch 11 train_loss=0.386 val_loss=0.402 val_f1=0.284
wide      epoch 12 train_loss=0.387 val_loss=0.402 val_f1=0.284
```

```
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
```

Validation summary:

	experiment	val_f1	val_pr_auc	val_recall_at_70_precision	threshold
0	wide	0.838869	0.836950		1.0 0.10
1	compact	0.838746	0.843062		1.0 0.10
2	baseline	0.838427	0.832994		1.0 0.15


```
[16]: # Plot training history for each experiment (train loss, val loss, F1 across epochs)
history_metrics = ["train_loss", "loss", "f1"]

for result in experiment_results:
    history_df = pd.DataFrame(result["history"]) # convert per-epoch metrics to DataFrame
    plt.figure(figsize=(7, 4))


    # Plot selected metrics if present
    for metric in history_metrics:
        if metric in history_df.columns:
            plt.plot(history_df["epoch"], history_df[metric], label=metric.replace("_", " "))


    plt.title(f"Training history - {result['name']}") # experiment name
    plt.xlabel("Epoch")
    plt.ylabel("Value")
    plt.grid(alpha=0.3)
    plt.legend()
    plt.tight_layout()
    plt.show()
```

Training history - baseline



Training history - compact





1.7 Test evaluation

Now we pick the best model from experiments (based on F1 score), evaluate the model and save the metrics, create classification report, confusion matrix and log it to Wandb

```
[17]: # Select best experiment by highest validation F1
best_experiment = max(experiment_results, key=lambda item: item["val_metrics"]["f1"])
best_cfg = best_experiment["config"] # best hyperparameters
best_threshold = best_experiment["threshold"] # tuned decision threshold (from val set)

# Recreate the best model architecture and load its saved weights
best_model = CovidMLP(len(feature_cols), best_cfg["hidden_layers"], best_cfg["dropout"]).to(device)
best_model.load_state_dict(best_experiment["state_dict"])

# Evaluate on the held-out test set
test_metrics_raw, test_probs, test_labels = evaluate_model(best_model, test_loader)

# Recompute metrics with the tuned threshold (keep ROC AUC & loss from raw eval)
test_metrics = metrics_from_probs(test_probs, test_labels, threshold=best_threshold)
test_metrics["roc_auc"] = test_metrics_raw["roc_auc"]
test_metrics["loss"] = test_metrics_raw["loss"]
```

```

# Pretty-print main test metrics
print("Test metrics with tuned threshold:")
for key in ["loss", "roc_auc", "pr_auc", "accuracy", "precision", "recall", "f1", "recall_at_70_precision"]:
    print(f" {key}: {test_metrics[key]:.4f}")

# Derive hard predictions using tuned threshold, then print a detailed classification report
binary_preds = (test_probs >= best_threshold).astype(int)
print("Classification report:")
print(classification_report(test_labels, binary_preds, digits=4, zero_division=0))

# Confusion matrix (as a DataFrame for nicer display)
cm = confusion_matrix(test_labels, binary_preds)
cm_df = pd.DataFrame(cm, index=["Actual 0", "Actual 1"], columns=["Pred 0", "Pred 1"])
print("Confusion matrix:")
display(cm_df)

# Log confusion matrix to Weights & Biases for the best config
wandb_run = wandb.init(project="covid-19", config=best_cfg, reinit=True) # new run for test artifacts
wandb_run.log({
    "confusion_matrix": wandb.plot.confusion_matrix(
        preds=binary_preds.tolist(),
        y_true=test_labels.tolist(),
        title="Confusion Matrix of best Model"
    )
})
wandb_run.finish()

# Plot Precision-Recall curve with the chosen operating point highlighted
precisions, recalls, _ = precision_recall_curve(test_labels, test_probs)
plt.figure(figsize=(6, 4))
plt.plot(recalls, precisions, label="PR curve")
plt.scatter(test_metrics["recall"], test_metrics["precision"], color="red", label="Chosen threshold")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Test precision-recall curve")
plt.grid(alpha=0.3)
plt.legend()
plt.show()

```

Test metrics with tuned threshold:

```
loss: 0.4024
roc_auc: 0.6921
pr_auc: 0.8346
accuracy: 0.7292
precision: 0.7301
recall: 0.9844
f1: 0.8384
recall_at_70_precision: 1.0000
```

Classification report:

	precision	recall	f1-score	support
0.0	0.7061	0.0932	0.1647	31470
1.0	0.7301	0.9844	0.8384	78396
accuracy			0.7292	109866
macro avg	0.7181	0.5388	0.5015	109866
weighted avg	0.7232	0.7292	0.6454	109866

Confusion matrix:

	Pred 0	Pred 1
Actual 0	2934	28536
Actual 1	1221	77175

```
<IPython.core.display.HTML object>
```

